

Homework 4 — Word embeddings

Overview

The large number of English words can make language-based applications daunting. To cope with this, it is helpful to have a *clustering* or *embedding* of these words, so that words with similar meanings are clustered together, or have embeddings that are close to one another.

But how can we get at the meanings of words? John Firth (1957) put it thus:

You shall know a word by the company it keeps.

That is, words that tend to appear in similar contexts are likely to be related. In this assignment, you will investigate this idea by coming up with an embedding of words that is based on co-occurrence statistics.

The description here assumes you are using Python with NLTK.

- First, download the Brown corpus (using `nltk.corpus`). This is a collection of text samples from a wide range of sources, with a total of over a million words. Calling `brown.words()` returns this text in one long list, which is useful.
- Remove stopwords and punctuation, make everything lowercase, and count how often each word occurs. Use this to come up with two lists:
 - A *vocabulary* V , consisting of a few thousand (e.g., 5000) of the most commonly-occurring words.
 - A shorter list C of at most 1000 of the most commonly-occurring words, which we shall call *context words*.
- For each word $w \in V$, and each occurrence of it in the text stream, look at the surrounding window of four words (two before, two after):

$w_1 \quad w_2 \quad w \quad w_3 \quad w_4.$

Keep count of how often context words from C appear in these positions around word w . That is, for $w \in V, c \in C$, define

$n(w, c) = \#$ of times c occurs in a window around w .

Using these counts, construct the probability distribution $\Pr(c|w)$ of context words around w (for each $w \in V$), as well as the overall distribution $\Pr(c)$ of context words. These are distributions over C .

- Represent each vocabulary item w by a $|C|$ -dimensional vector $\Phi(w)$, whose c 'th coordinate is:

$$\Phi_c(w) = \max \left(0, \log \frac{\Pr(c|w)}{\Pr(c)} \right).$$

This is known as the (positive) *pointwise mutual information*, and has been quite successful in work on word embeddings.

- Suppose we want a 100-dimensional representation. How would you achieve this?

- Investigate the resulting embedding in two ways:
 - Cluster the vocabulary into 100 clusters. Look them over; do they seem completely random, or is there some sense to them?
 - Try finding the nearest neighbor of selected words. Do the answers make sense?

The Brown corpus is very small. Current work on word embeddings uses data sets that are several orders of magnitude larger, but the methodology is along the same lines.

What to turn in

On the due date, turn in a **typewritten** report containing the following elements (each labeled clearly).

1. *A description of your 100-dimensional embedding.*

The description should be concise and clear, and should make it obvious exactly what steps you took to obtain your word embeddings. Below, we will denote these as $\Psi(w) \in \mathbb{R}^{100}$, for $w \in V$. Also clarify exactly how you selected the vocabulary V and the context words C .

2. *Nearest neighbor results.*

Pick a collection of 25 words $w \in V$. For each w , return its nearest neighbor $w' \neq w$ in V . A popular distance measure to use for this is *cosine distance*:

$$1 - \frac{\Psi(w) \cdot \Psi(w')}{\|\Psi(w)\| \|\Psi(w')\|}.$$

Here are some suggestions for words you might choose:

communism, autumn, cigarette, pulmonary, mankind, africa, chicago, revolution, september,
chemical, detergent, dictionary, storm, worship

Do the results make any sense?

3. *Clustering.*

Using the vectorial representation $\Psi(\cdot)$, cluster the words in V into 100 groups. Clearly specify what algorithm and distance function you using for this, and the reasons for your choices.

Look over the resulting 100 clusters. Do any of them seem even moderately coherent? Pick out a few of the best clusters and list the words in them.